

A Scalable Algorithm for Maximizing Range Sum in Spatial Databases


Dong-Wan Choi, KAIST, Korea

Chin-Wan Chung, KAIST, Korea

Yufei Tao, KAIST & CUHK, Hong Kong


MOTIVATION

- ▶ Where is **the most representative spot** in a city for a tourist?



- ▶ Maximize the number of tourist attractions within a movable range





MOTIVATION



- ▶ Where is **the most profitable place** to set up a new pizza store?
 - ▶ Maximize the number of residents within a delivery range

MOTIVATION

- ▶ Where is **the best location** that maximizes the number of objects covered by **a given range**?



FORMAL DEFINITION

- ▶ MaxRS (Maximizing Range Sum) problem
 - ▶ Given a set O of weighted objects and a rectangle r of a given size,
 - ▶ Find a location p of r that maximizes:
$$\sum_{o \in O_{r(p)}} w(o),$$
 where
 $r(p)$ is the rectangle centered at a location p ,
 $O_{r(p)}$ is the set of objects covered by $r(p)$, and
 $w(o)$ is the weight of $o \in O$
- ▶ MaxCRS (Maximizing Circular Range Sum) problem
 - ▶ The circle version of the MaxRS problem



CONTENTS

- ❖ ~~Motivation~~
- ❖ ~~Formal Definition~~
- ❖ Preliminary
- ❖ Our Algorithms
 - ❖ ExactMaxRS
 - ❖ ApproxMaxCRS
- ❖ Theoretical Results
- ❖ Experimental Results
- ❖ Conclusion and Future works

PRELIMINARY

- ▶ A naïve solution
 - ▶ Issuing range aggregate queries for every location
 - ▶ **Problem: Infinite # of locations!**

- ▶ Problem transformation




- ▶ Given a set of rectangles, find **the most dense region** where the most rectangles intersect

EXACTMAXRS

- ▶ Exact algorithm that solves MaxRS
- ▶ External-memory algorithm
 - ▶ Scalable for a massive dataset
- ▶ Follows the divide-and-conquer strategy:
 - ▶ **Recursively divide** the entire dataset into smaller subsets
 - ▶ **Compute a local solution** for each subset
 - ▶ **Merge** local solutions of subsets


DIVISION PHASE



- ▶ **Divide the space vertically** into m sub-spaces, called slabs, each of which has roughly the same # of rectangles
 - ▶ Until the # of rectangles can fit in the main memory
- ▶ **Do not pass spanning rectangles** to the next level of recursion

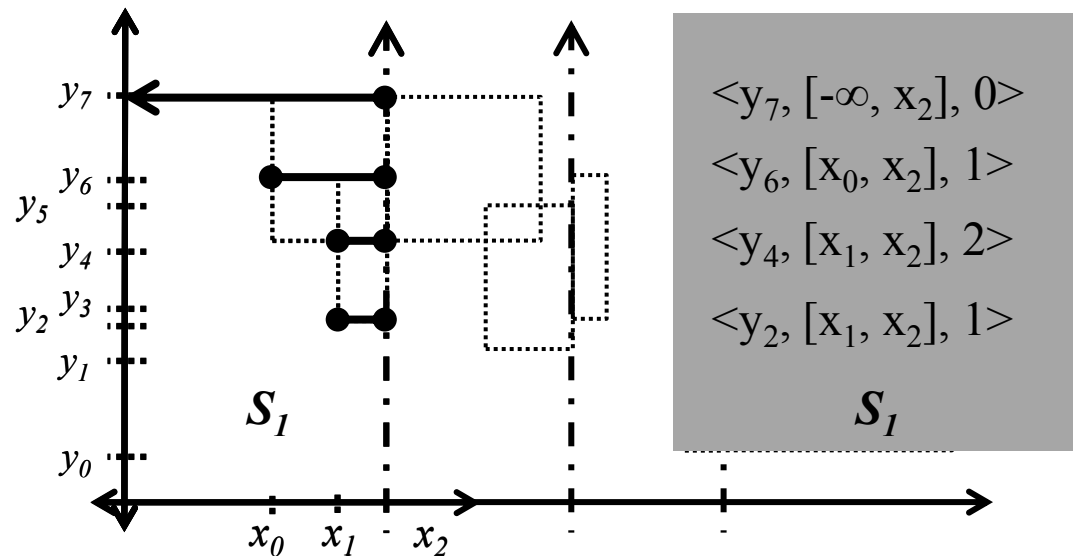
SLAB-FILES

- ▶ The structure to be returned **after conquering the sub-problem** w.r.t. a slab
- ▶ The set of tuples, each of which is $t = \langle y, [x_i, x_j], sum \rangle$
 - ▶ In the upward direction, after y , the most dense region (whose total weights is sum) is in the x-range $[x_i, x_j]$.




SLAB-FILES

► An example of a slab-file



MERGING PHASE





- ▶ Sweep a horizontal line across the slab-files (S_1, \dots, S_4) and the spanning rectangle file (R')
- ▶ When encountering several tuples at a horizontal line, choose a tuple with a maximum sum

APPROXMAXCRS

- ▶ Approximation algorithm for MaxCRS
- ▶ Uses the ExactMaxRS algorithm as a tool
- ▶ Overall Flow
 1. **Transform** MaxCRS into MaxRS with **MBRs**
 2. **Do ExactMaxRS** on the transformed dataset
 3. Generate **candidate points** based on the result from ExactMaxRS
 4. **Choose** the best point among the candidate points

TRANSFORMATION


- ▶ MaxCRS \rightarrow MaxRS with MBRs
 - ▶ Construct the MBR for each circle



- ▶ Find the most dense region of MaxRS with MBRs

CANDIDATE POINTS

- ▶ Generate **five candidate points** based on the result from ExactMaxRS
 - ▶ p_0 : **the center point** of the most dense region returned from ExactMaxRS
 - ▶ p_1, p_2, p_3, p_4 : **four shifted points** from p_0



$$\sigma \in \left((\sqrt{2} - 1) \frac{d}{2}, \frac{d}{2} \right)$$

, where d is the diameter of circles

- ▶ **Return the best point p_i** among p_0, \dots, p_4
 - ▶ such that the total weight of the circles covering p_i is maximized.

THEORETICAL RESULTS

▶ ExactMaxRS

- ▶ **Optimal** in terms of the I/O complexity
 - ▶ $O((N/B)\log_{M/B}(N/B))$ I/O's, where N is the # of objects, M is the memory size, and B is the block size
 - The counterpart of $O(n\log n)$ in the main memory
→ optimal time complexity

▶ ApproxMaxCRS

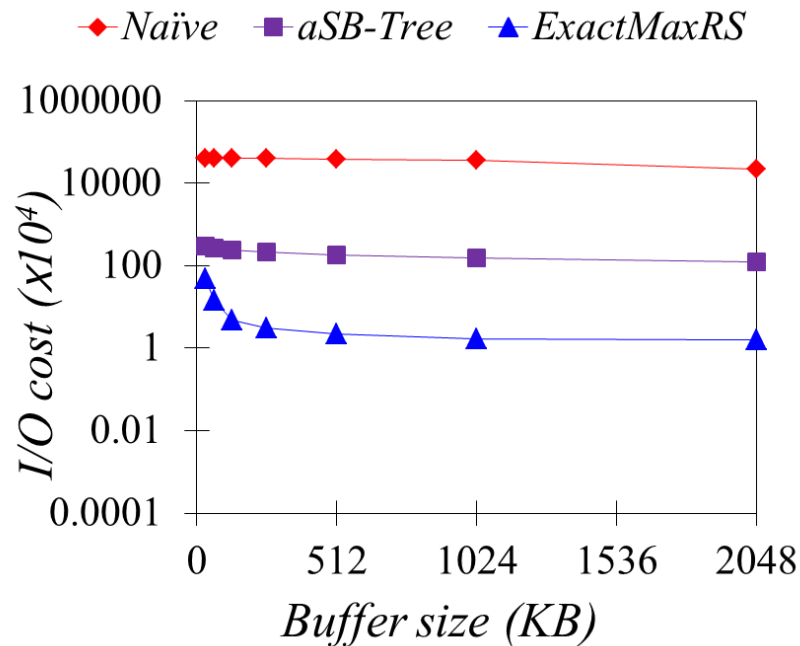
- ▶ **(1/4)-approximation bound**
 - ▶ Much better in practice

EXPERIMENTAL RESULTS


- ▶ Environment Setting
 - ▶ Synthetic datasets
 - ▶ Uniform distribution and Gaussian distribution
 - ▶ Cardinality: 100,000~500,000
 - ▶ Real datasets
 - ▶ NE and UX datasets from the R-tree Portal site
 - ▶ Cardinality: 123,593 for NE, 19,499 for UX
 - ▶ Compared approaches
 - ▶ Naïve plane sweep, $O(n^2)$
 - ▶ Plane sweep using aSB-tree, $O(n \log n)$
 - Adaptation of the optimal in-memory algorithm for the MaxRS problem

EXPERIMENTAL RESULTS

► I/O cost with varying the buffer size




<Gaussian distribution>



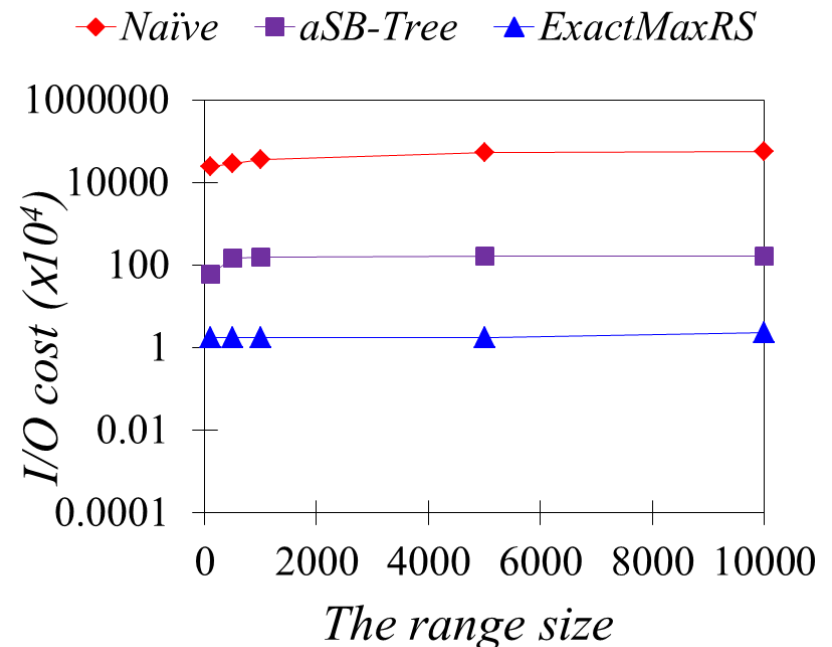
<Uniform distribution>

EXPERIMENTAL RESULTS

► I/O cost with varying the range size



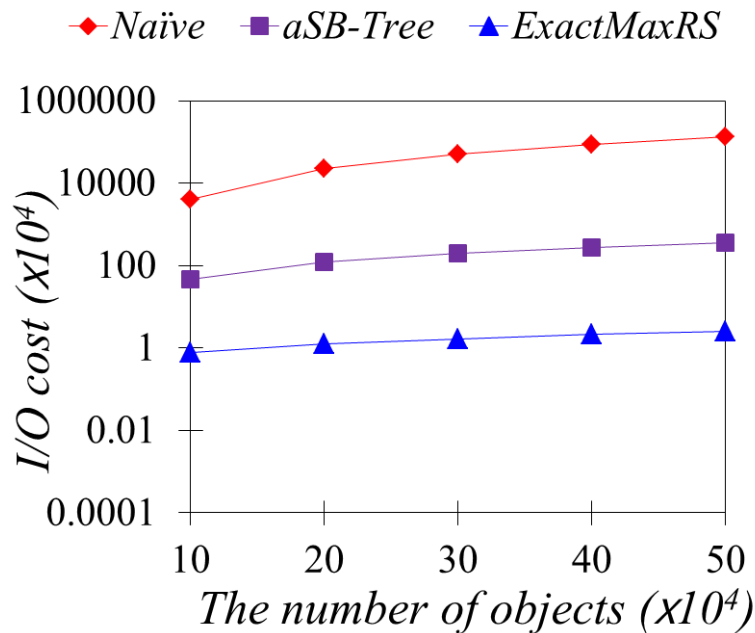
<Gaussian distribution>




<Uniform distribution>

EXPERIMENTAL RESULTS

► I/O cost with varying the cardinality



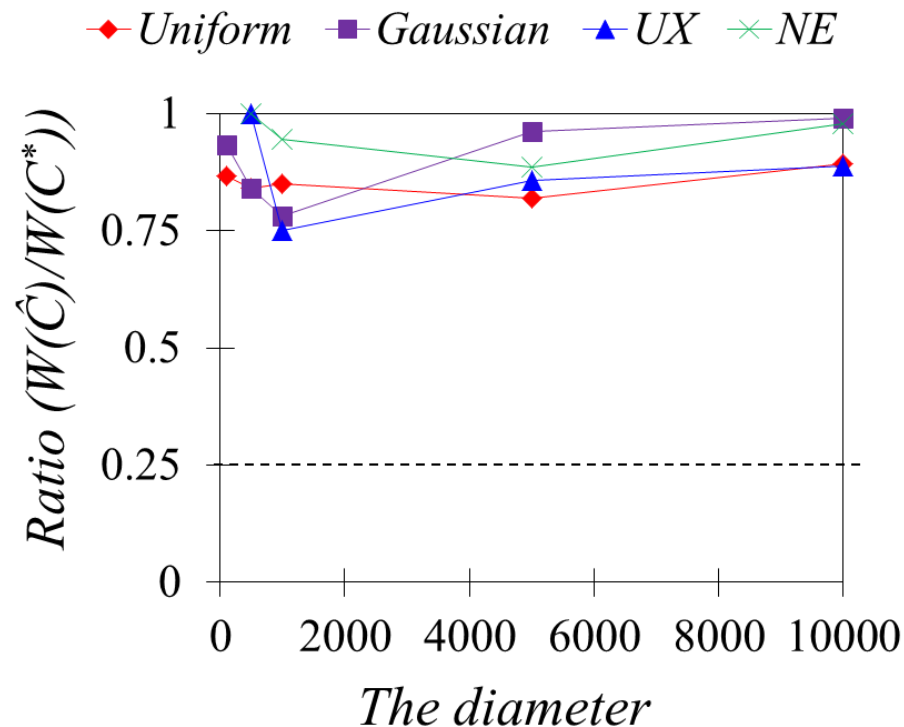
<Gaussian distribution>



<Uniform distribution>

EXPERIMENTAL RESULTS

- ▶ Error ratio of ApproxMaxCRS with varying the diameter



CONCLUSION & FUTURE WORKS

▶ Contributions

- ▶ Proposed **the first optimal external-memory algorithm** for the MaxRS problem
- ▶ Proposed **the (1/4)-approximation algorithm** for the MaxCRS problem
- ▶ Proved theoretically the **correctness, optimality, approximation bound**, and **tightness** of the bound
- ▶ Evaluated experimentally the efficiency and accuracy

▶ Future works

- ▶ Max k RS problem
- ▶ MinRS problem



Thank You

APPENDIX

- ▶ The proof sketch of the correctness of ExactMaxRS
 - ▶ The x-range of a tuple in a slab-file is called "max-interval".
 - ▶ Let I^* be the max-interval w.r.t. entire space and I^\wedge be a piece of I for a recursion level. Then I^\wedge is also the max-interval w.r.t. the sub-space corresponding to the recursion level.
 - ▶ There cannot be an interval I' whose sum is larger than I^\wedge , since spanning rectangles can affect all or none of intervals in a slab.


APPENDIX

- ▶ The proof sketch of $(1/4)$ -approximation bound
- ▶ The four circles together cover the region for MaxRS, which covers k points. Hence, at least one of those circles covers $k/4$ points.



APPENDIX

- ▶ The proof sketch of the tightness of the approximation bound
 - ▶ $4p_i \leq p^*$



APPENDIX

► Results on real datasets

